

## Parallel Computing in R




# Introduction to Parallelization

# Goals for this Lecture


- Review pros and cons of parallel/distributed computing
- Introduce parallel computing methods in R
- Methods that parallelize general loops
  - The `parallel` package and `mclapply` commands.
  - The `foreach` and `doParallel` packages.
- Methods that are `dplyr`-aware
  - The `multidplyr` package
  - Local Spark instances with `sparklyr`
- Types of problems that are amenable to each approach
- Effects of adding cores/clusters

# R is single-threaded by design, but there are exceptions.

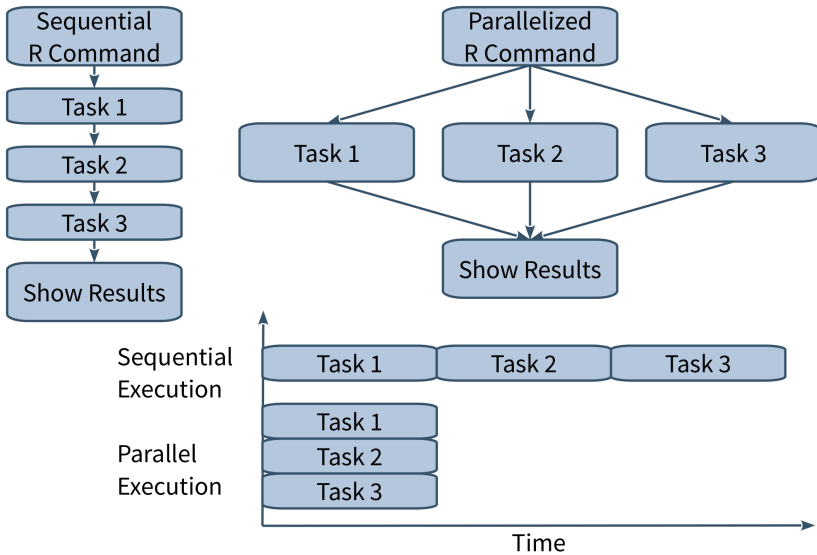
Single-Threaded R Code



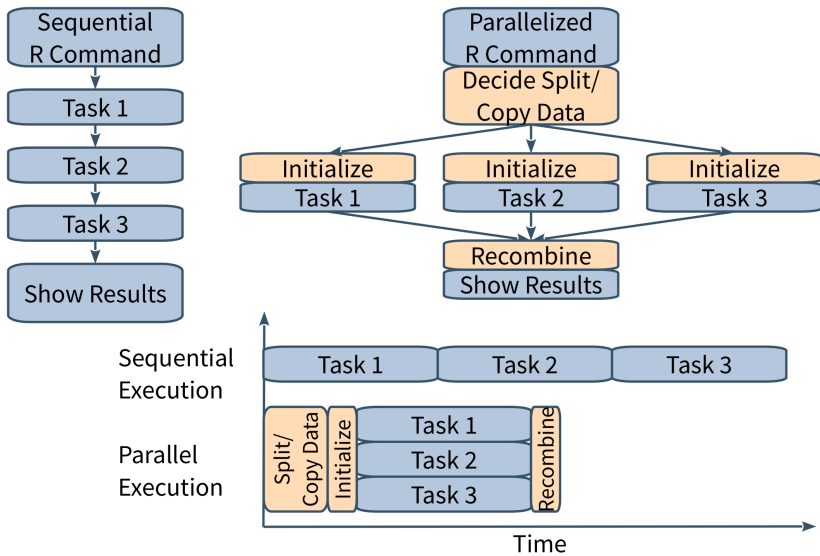
Some R Versions/Commands Are Multi-Threaded




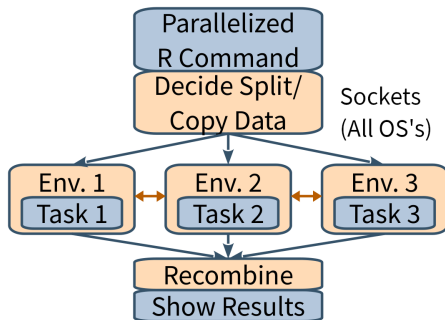
In a perfect world, parallelization would be 100% efficient.



In reality, there is always overhead in splitting a task. How much depends on the nature of the task.



Efficiency is also affected by whether or not processes share memory, and by how they must communicate.



- Separate environments are more costly in time and resources.
- Communication between tasks adds complexity.

## Task assignment can be pre-scheduled or dynamic.

- Pre-scheduling breaks the tasks into chunks equal to the number of cores/threads/cluster nodes at the beginning.
  - Good for short, similar multitudinous tasks.
  - Good when task length can be easily estimated for balancing.
- Dynamic assignment sends the next task to the next available core/thread/node.
  - Good when task time is highly variable.
  - Good for smaller numbers of tasks.
  - More overhead?





## Simulation: Parallelizing Loops

# Hypothetical Example: Testing a weight-loss program.

- You want to test a new weight loss program over one month.
- *If the program works*, it would be reasonable for subjects to lose five pounds, on average.
- Assume two groups with a real difference in means:
  - Control group. Mean weight loss = 0, std. dev. = 3 lbs.
  - Treatment group. Mean weight loss = 5, std. dev. = 3 lbs.
- Will two groups of size  $n$  be able to detect that difference?

```
diff.test <- function(n=5, a=0.05){  
  t.test(x=rnorm(n, mean=0, sd=3), y=rnorm(n, mean=5, sd=3))$p.value < a  
}  
diff.test(n=5, a=0.05)
```

```
## [1] FALSE
```

Statistical *power* is the probability that a hypothesis test will detect a true difference between groups.

- Power *can* be calculated theoretically in some cases.
- Simulation can be used to calculate power even when theory is difficult.
  - Assume a certain probability distribution for each group.
  - Simulate a random sample from each group's distribution.
  - Do the test, and record whether the test detects a difference (rejects  $H_0$ ). [This is the `test.diff` function.]
  - Repeat.
  - Calculate the percentage of runs that detect a difference.

# Without parallelization, we could implement the simulation in *many* ways!

```
# Replicate repeats expressions  
# that don't need inputs.  
test.vec <- replicate(10000, diff.test())  
mean(test.vec)
```

```
## [1] 0.6064
```

```
# lapply plugs each value of n.vec into  
# the function and returns a list.
```

```
n.vec <- rep(5, 10000)  
test.vec <- lapply(n.vec, diff.test)  
mean(unlist(test.vec))
```

```
## [1] 0.61
```

```
# A standard `for` loop to do the same
```

```
test.vec <- numeric(10000)  
for(i in 1:10000){  
  test.vec[i] <- diff.test()  
}  
mean(test.vec)
```

```
## [1] 0.608
```

```
# `foreach` gives a new syntax, which acts  
# like a `for` loop, but combines output.
```

```
library(foreach)  
test.vec <-  
  foreach(n=n.vec, .combine=c) %do%  
    diff.test(n)  
mean(test.vec)
```

```
## [1] 0.6237
```

```
# The `map` family from `purrr` are similar  
# to `lapply` for predefined functions.
```

```
test.vec <- map_dbl(n.vec, diff.test)  
# but with a nice function construction  
# syntax.
```

```
test.vec <- map_dbl(n.vec,  
  ~t.test(x=rnorm(.x, mean=0, sd=3),  
    y=rnorm(.x, mean=5, sd=3))$p.value)  
mean(test.vec)
```

```
## [1] 0.6091
```

# the parallel package provides mclapply which uses the standard lapply syntax. (linux/unix/mac)

## Single-thread computation

```
# want 10000 repetitions, and  
# lapply needs an input.
```

```
n.vec <- rep(5, 10000)  
system.time({  
  test.vec <-  
    lapply(n.vec, diff.test)  
})
```

```
##      user  system elapsed  
##    1.324    0.000    1.324
```

```
test.power <-  
  mean(unlist(test.vec))  
test.power  
## [1] 0.6111
```

## Multi-thread computation

```
# note: mclapply requires forks.  
# needs linux/mac (not windows).
```

```
library(parallel)  
n.cores <- detectCores()-1  
system.time({  
  test.vec <- mclapply(n.vec,  
    diff.test,  
    mc.cores=n.cores)  
})
```

```
##      user  system elapsed  
##    2.102    0.341    0.519
```

```
test.power <-  
  mean(unlist(test.vec))  
test.power  
## [1] 0.6078
```

# The parallel package also provides parLapply which is used with pre-defined clusters (all platforms).

## Timing Excluding Initialization

```
cl <- makeCluster(n.cores)
system.time({
  test.vec <- parLapply(cl, n.vec,
    diff.test)
})

##      user  system elapsed
##    0.001   0.004   0.402

# Make sure to stop your cluster!
stopCluster(cl)
test.power <-
  mean(unlist(test.vec))
test.power

## [1] 0.6163
```

## Initialization adds Overhead

```
system.time({
  cl <- makeCluster(n.cores)
  test.vec <- parLapply(cl, n.vec,
    diff.test)
# Make sure to stop your cluster!
stopCluster(cl)
})

##      user  system elapsed
##    0.010   0.000   0.983

test.power <-
  mean(unlist(test.vec))
test.power

## [1] 0.6131
```

# The doParallel package provides an parallel backend to the foreach package with platform-independent syntax.

## Single-Thread Computation

```
library(foreach)
system.time({
  test.vec <- foreach(n=n.vec) %do%
    diff.test(n)
})

##      user      system elapsed
##    3.221      0.000      3.220

test.power <-
  mean(unlist(test.vec))
test.power

## [1] 0.6137
```

## Multi-Thread Computation

```
library(doParallel)
registerDoParallel(n.cores)
system.time({
  test.vec <-
    foreach(n=n.vec, .combine=c) %dopar%
      diff.test(n)
})

##      user      system elapsed
##    3.973      0.412      1.598

test.power <- mean(test.vec)
test.power

## [1] 0.6107
```

**Note:** `registerDoParallel(<number>)` selects the correct method for your system, and automatically cleans up clusters when done.

When cluster initialization is included, forks (Unix/Linux) appear faster than PSOCK clusters (all platforms).

## Forks

```
system.time({  
  cl <-  
    makeCluster(n.cores, type="FORK")  
  registerDoParallel(cl)  
  test.vec <-  
    foreach(n=n.vec) %dopar%  
      diff.test(n)  
})  
  
##      user  system elapsed  
##    2.519   0.276   2.820  
  
stopCluster(cl)  
test.power <- mean(unlist(test.vec))  
test.power  
  
## [1] 0.6167
```

## PSOCK Clusters


```
system.time({  
  cl <-  
    makeCluster(n.cores, type="PSOCK")  
  registerDoParallel(cl)  
  test.vec <-  
    foreach(n=n.vec) %dopar%  
      diff.test(n)  
})  
  
##      user  system elapsed  
##    2.596   0.216   3.324  
  
stopCluster(cl)  
test.power <- mean(unlist(test.vec))  
test.power  
  
## [1] 0.6152
```

**Note:** If clusters are defined explicitly, they must be stopped as well.




Because of additional overhead, completion time does not decrease much past a certain number of cores.

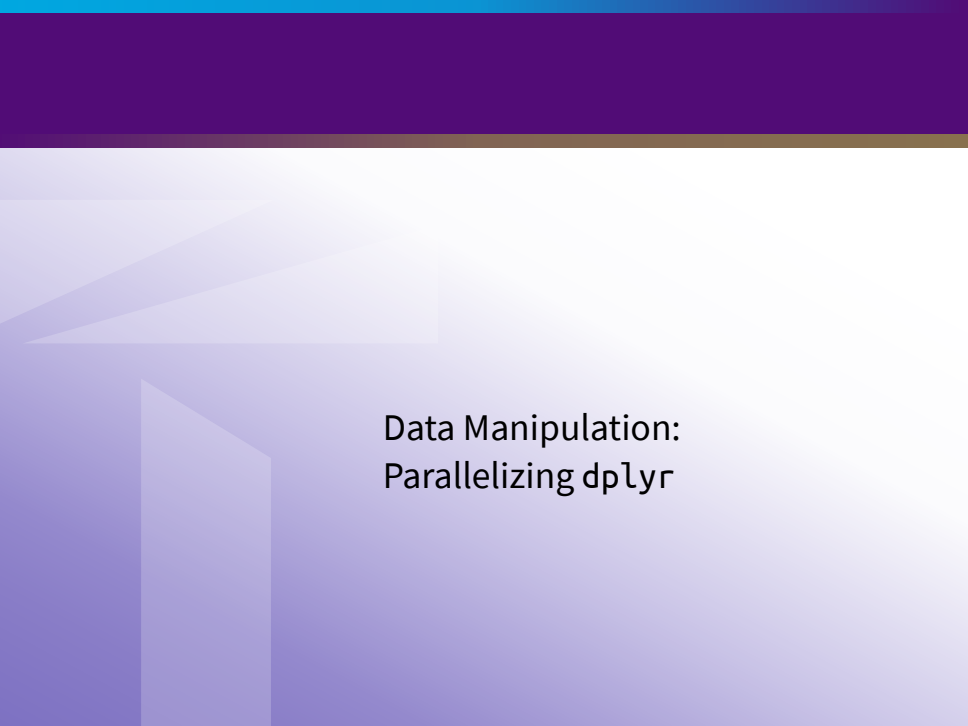
```
# apply loops through numbers of cores and creates a vector of results.
time.vec <- sapply(1:n.cores, function(cores){
  system.time{
    test.vec <- mclapply(n.vec, diff.test, mc.cores=cores)
  }["elapsed"]
})
data.frame(n.cores = 1:n.cores, Actual = time.vec,
  Ideal = system.time(lapply(n.vec, diff.test))["elapsed"]/(1:n.cores)) %>%
pivot_longer(-n.cores, names_to="Type", values_to="Time") %>%
ggplot(aes(x=n.cores, y=Time, color=Type)) + geom_point() + geom_line()
```



Because of additional overhead, completion time does not decrease much past a certain number of cores.

```
time.vec <- apply(1:n.cores, function(cores){  
  registerDoParallel(cores)  
  system.time{  
    test.vec <- foreach(n=n.vec, .combine=c) %dopar% diff.test(n)  
  }["elapsed"]  
})  
ideal.t <- system.time(foreach(n=n.vec) %do% diff.test(n))["elapsed"]  
data.frame(n.cores = 1:n.cores, Actual = time.vec, Ideal=ideal.t/(1:n.cores))  
  pivot_longer(-n.cores, names_to="Type", values_to="Time") %>%  
  ggplot(aes(x=n.cores, y=Time, color=Type)) + geom_point() + geom_line()
```





## Data Manipulation: Parallelizing dplyr

## Example: Iris Averages

What are the average values for the parameters measured in the iris data set?

```
head(iris)
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 1	5.1	3.5	1.4	0.2	setosa
## 2	4.9	3.0	1.4	0.2	setosa
## 3	4.7	3.2	1.3	0.2	setosa
## 4	4.6	3.1	1.5	0.2	setosa
## 5	5.0	3.6	1.4	0.2	setosa
## 6	5.4	3.9	1.7	0.4	setosa

# A Single-Core Solution

```
my.iris <- iris %>%  
  rename(SL=Sepal.Length, SW=Sepal.Width,  
         PL=Petal.Length, PW=Petal.Width)  
my.iris %>% group_by(Species) %>% summarize(across(everything(), mean)  
)
```

```
## # A tibble: 3 x 5  
##   Species      SL      SW      PL      PW  
##   <fct>      <dbl> <dbl> <dbl> <dbl>  
## 1 setosa      5.01  3.43  1.46 0.246  
## 2 versicolor 5.94  2.77  4.26 1.33  
## 3 virginica  6.59  2.97  5.55 2.03
```

# The `multidplyr` package creates a new `dplyr` backend that uses parallel processing.

Existing `dplyr` code can be used, with only a few additions:

- `cl <- new_cluster(<number>)` defines a cluster with `<number>` cores.
- Each cluster runs a new R instance, and must be initialized.
  - `cluster_library(cl, <packages>)` loads libraries from character vector `<packages>`.
  - `cluster_copy(cl, <names>)` copies variables from character vector `<names>`.
- `partition(cl)` assigns data frame groups to various clusters.
- `collect()` combines results from each subprocess.

# Parallelizing doesn't improve run time on small data sets. Data preparation overhead is significant.

```
system.time({  
  my.iris %>% group_by(Species) %>%  
  summarize(across(everything(), mean)) %>%  
  print()  
})
```

```
## # A tibble: 3 x 5  
##   Species      SL      SW      PL      PW  
##   <fct>      <dbl> <dbl> <dbl> <dbl>  
## 1 setosa      5.01  3.43  1.46  0.246  
## 2 versicolor 5.94  2.77  4.26  1.33  
## 3 virginica  6.59  2.97  5.55  2.03  
##   user system elapsed  
##  0.024  0.000  0.023
```

```
library(multidplyr)  
cl <- new_cluster(7)  
system.time({  
  my.iris %>% group_by(Species) %>%  
  partition(cl) %>%  
  summarize(across(everything(), mean)) %>%  
  collect() %>%  
  print()  
})
```

```
## Using partial cluster of size 3  
## # A tibble: 3 x 5  
##   Species      SL      SW      PL      PW  
##   <fct>      <dbl> <dbl> <dbl> <dbl>  
## 1 setosa      5.01  3.43  1.46  0.246  
## 2 versicolor 5.94  2.77  4.26  1.33  
## 3 virginica  6.59  2.97  5.55  2.03  
##   user system elapsed  
##  0.073  0.004  0.994
```

# Timing that doesn't include data/cluster preparation shows some improvement.

```
cl <- new_cluster(7)
system.time({
  my.iris.grp <- my.iris %>%
    group_by(Species) %>% partition(cl)
  my.iris.grp %>%
    summarize(across(everything(), mean)) %>%
    collect() %>%
    print()
})
```

```
## Using partial cluster of size 3
## # A tibble: 3 x 5
##   Species      SL      SW      PL      PW
##   <fct>      <dbl> <dbl> <dbl> <dbl>
## 1 setosa      5.01  3.43  1.46  0.246
## 2 versicolor 5.94  2.77  4.26  1.33
## 3 virginica  6.59  2.97  5.55  2.03
##   user system elapsed
## 0.057  0.001  0.943
```

```
cl <- new_cluster(7)
my.iris.grp <- my.iris %>%
  group_by(Species) %>% partition(cl)
## Using partial cluster of size 3
system.time({
  my.iris.grp %>%
    summarize(across(everything(), mean)) %>%
    collect() %>%
    print()
})
```

```
## # A tibble: 3 x 5
##   Species      SL      SW      PL      PW
##   <fct>      <dbl> <dbl> <dbl> <dbl>
## 1 setosa      5.01  3.43  1.46  0.246
## 2 versicolor 5.94  2.77  4.26  1.33
## 3 virginica  6.59  2.97  5.55  2.03
##   user system elapsed
## 0.059  0.000  0.919
```





## Parallel Computation with sparklyr

First, install the `sparklyr` package and a local version of Spark.

```
install.packages("sparklyr")  
spark_install()
```

- Note that `spark_install` options can specify the Spark version and other options, but the defaults usually work.
- By default, Spark will then create nodes on your local machine which will use multiple cores.

# The sparklyr package allows use of dplyr syntax to invoke Spark.

- `spark_connect()` connects to a spark session. In this case, the local session that's installed by default.
- `copy_to()` prepares data for use with spark.
- Always close your connection with `spark_disconnect()!`

```
library(sparklyr)
sc <- spark_connect(master = "local")
system.time({
  my.iris.tbl <- copy_to(sc, my.iris)
  my.iris.tbl %>%
    group_by(Species) %>%
    # Note: everything() doesn't work here.
    summarize(SL=mean(SL ), SW=mean(SW)) %>%
    print()
})

## # Source: spark<?> [?? x 3]
##   Species      SL    SW
##   <chr>        <dbl> <dbl>
## 1 versicolor  5.94  2.77
## 2 virginica   6.59  2.97
## 3 setosa      5.01  3.43
##   user  system elapsed
##  0.376   0.020   8.140

spark_disconnect(sc)
```

# Data preparation constitutes a large chunk of this small problem's run time.

```
sc <- spark_connect(master = "local")
system.time({
my.iris.tbl <- copy_to(sc, my.iris)
my.iris.tbl %>%
  group_by(Species) %>%
  # Note: everything() doesn't work here.
  summarize(SL=mean(SL ), SW=mean(SW)) %>%
  print()
})
```

```
## # Source: spark<?> [?? x 3]
##   Species      SL    SW
##   <chr>        <dbl> <dbl>
## 1 versicolor  5.94  2.77
## 2 virginica   6.59  2.97
## 3 setosa      5.01  3.43
##   user  system elapsed
## 0.297  0.000  7.853
```

```
spark_disconnect(sc)
```

```
sc <- spark_connect(master = "local")
my.iris.tbl <- copy_to(sc, my.iris)
system.time({
my.iris.tbl %>%
  group_by(Species) %>%
  # Note: everything() doesn't work here.
  summarize(SL=mean(SL ), SW=mean(SW)) %>%
  print()
})
```

```
## # Source: spark<?> [?? x 3]
##   Species      SL    SW
##   <chr>        <dbl> <dbl>
## 1 versicolor  5.94  2.77
## 2 virginica   6.59  2.97
## 3 setosa      5.01  3.43
##   user  system elapsed
## 0.133  0.000  1.339
```

```
spark_disconnect(sc)
```



## Parallelization Summary

# Summary of Parallelization Methods

Sequential	Parallelized	Type	Package
<-	%<-%	Assignment	future
map family	future_map family	Looping	furrr
lapply	mclapply/parLapply	Looping	parallel
foreach() %do%	foreach %dopar%	Looping	doParallel
dplyr cmds	dplyr cmds	Data Analysis	multidplyr
dplyr cmds	dplyr cmds	Data Analysis	sparklyr