

PDAT615G: Machine Learning

Module 1 – Introduction

Module 1: Introduction

- Supervised, unsupervised and reinforcement learning
- Regression models
 - Single/multiple variable for numeric prediction
 - Logistic regression for classification
- Review of R programming techniques
 - Defining functions
 - Vectorized commands & loop-like constructions



Types of Machine Learning

Supervised learning trains a model to use certain features, or variables, to predict a response.

- Training data contains both explanatory variables and the response.
- Models are trained to be able to use only the explanatory variables to predict new responses.
- Regression is an example of a supervised learning technique.

Unsupervised learning analyzes patterns in data without a particular response variable specified.

- Clustering is one example. What groupings appear in the data?
- Dimension reduction is another. Do fewer variables describe the data just as well?

Reinforcement learning creates an algorithm by rewarding “good” behavior and punishing “bad” behavior.

- Reinforcement learners need
 - a way to measure the desirability of the end state,
 - methods to explore the space of all possible actions, and
 - methods to find optimal strategies.
- Examples might include
 - a game-playing algorithm, or
 - a planner for supply-chain management.

Regression Models

Linear regression models can be used to predict a numeric response variable.

Linear regression models predict the value of one variable (the response, or dependent variable) based on other variables (the predictor, or independent, variables):

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_k x_k + \varepsilon.$$

- **Response (dependent) variable:** denoted Y

The response variable Y must be a numeric variable.

- **Explanatory (independent) variables:** denoted x_1, x_2, \dots, x_k

Predictors x_1, \dots, x_k can be numeric or categorical variables.

If we have **MORE THAN ONE** independent variable, say x_1 and x_2 , or x_1, x_2 and x_3 , the model is a *multiple linear regression model*.

Regression coefficients are chosen to minimize the sum of squared residuals.

- We assume a model form:

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_k x_k + \varepsilon.$$

- The β 's are estimated by b_0, b_1, \dots , giving the regression equation to calculate \hat{y} , the predicted value of y .

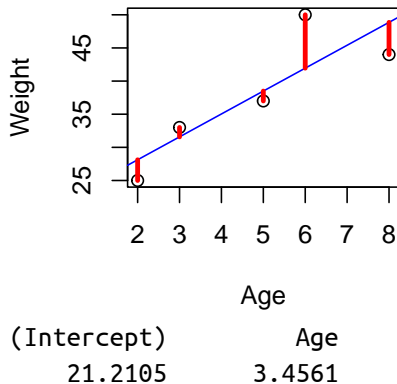
$$\hat{y} = b_0 + b_1 x_1 + b_2 x_2 + \cdots + b_k x_k.$$

- The coefficients b_0, b_1, \dots, b_k are chosen to minimize

$$\sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Regression models are calculated with `lm`, and predicted values obtained with `predict`.

```
df <- data.frame(  
  Age = c(2, 5, 3, 8, 6),  
  Weight = c(25, 37, 33, 44, 50)  
)  
wt.fit <- lm(Weight ~ Age, data=df)  
df$Pred.Wt <- predict(wt.fit)  
with(df, {  
  plot(x=Age, y=Weight)  
  abline(wt.fit, col="blue")  
  segments(x0=Age, x1=Age,  
           y0=Weight, y1=Pred.Wt,  
           col="red", lwd=3)  
})  
coef(wt.fit)
```



Logistic regression predicts the likelihood of a Yes/No categorical variable.


- Suppose p represents the probability of a “Yes” response.
- Perhaps p is related to one or more predictor variables.
- Logistic regression models the *log odds* as a linear function of one or more variables:

$$\log \left(\frac{p}{1-p} \right) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \varepsilon, \text{ or } p = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \cdots + \varepsilon)}}$$

- Note: Functions of the form $p = \frac{1}{1+e^{-z}}$ asymptotically approach 0 and 1 in an “S”-shaped curve, which is good for modeling probabilities.

Logistic regression models are calculated with the `glm` command and `family="binomial"`.

```
# Note: glm accepts 0/1 or
# factor variables.
df$Vax <- c(0, 1, 0, 1, 0)
vax.fit <- glm(Vax ~ Age, data=df,
               family="binomial")
with(df, plot(x=Age, y=Vax))
# Plot the curve
df.curve <- data.frame(
  Age = seq(2, 8, by=0.01)
)
df.curve$p <- predict(vax.fit,
                      newdata=df.curve, type="response")
with(df.curve,
     lines(x=Age, y=p, col="blue"))
```

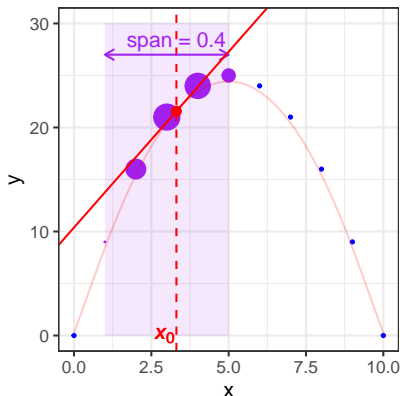


LOESS creates a “local” regression model near each point by using only near points in model construction.

- Historically, two related terms have been used.
- LOWESS – LOcally Weighted Scatterplot Smoothing
 - `lowess()`: Calculates a fit for one x and one y . Outputs predicted points for plotting.
 - Defaults to first-degree (linear) fitting.
- LOESS – LOcally Estimated Scatterplot Smoothing
 - Interface similar to `lm()`.
 - Outputs a model for use with `predict()`.
 - Generalizes LOWESS. Can accommodate multiple predictor variables.
 - Defaults to second-degree polynomial fitting.


Example: Degree 1 LOESS Fit, 40% Span

- We want to find a function $\hat{f}(x)$ that fits the data.
- Step 1: For a base point x_0 , find the closest 40% of the data points.
- Step 2: Weight each point within the span according to its distance from x_0 .
- Step 3: Calculate a weighted linear regression using those points.
- Step 4: $\hat{f}(x_0)$ is the point on that line above x_0
- Step 5: Repeat for other x 's.




Example: Degree 1 LOESS Fit, 40% Span

- We want to find a function $\hat{f}(x)$ that fits the data.
- Step 1: For a base point x_0 , find the closest 40% of the data points.
- Step 2: Weight each point within the span according to its distance from x_0 .
- Step 3: Calculate a weighted linear regression using those points.
- Step 4: $\hat{f}(x_0)$ is the point on that line above x_0
- Step 5: Repeat for other x 's.



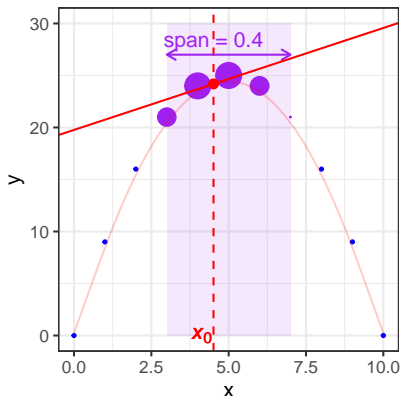
Example: Degree 1 LOESS Fit, 40% Span

- We want to find a function $\hat{f}(x)$ that fits the data.
- Step 1: For a base point x_0 , find the closest 40% of the data points.
- Step 2: Weight each point within the span according to its distance from x_0 .
- Step 3: Calculate a weighted linear regression using those points.
- Step 4: $\hat{f}(x_0)$ is the point on that line above x_0
- Step 5: Repeat for other x 's.




Example: Degree 1 LOESS Fit, 40% Span

- We want to find a function $\hat{f}(x)$ that fits the data.
- Step 1: For a base point x_0 , find the closest 40% of the data points.
- Step 2: Weight each point within the span according to its distance from x_0 .
- Step 3: Calculate a weighted linear regression using those points.
- Step 4: $\hat{f}(x_0)$ is the point on that line above x_0
- Step 5: Repeat for other x 's.




Example: Degree 1 LOESS Fit, 40% Span

- We want to find a function $\hat{f}(x)$ that fits the data.
- Step 1: For a base point x_0 , find the closest 40% of the data points.
- Step 2: Weight each point within the span according to its distance from x_0 .
- Step 3: Calculate a weighted linear regression using those points.
- Step 4: $\hat{f}(x_0)$ is the point on that line above x_0
- Step 5: Repeat for other x 's.



Increasing the span moves from a more “local” to a more “global” fit.



Note: Increasing span larger than 1 considers all the data points and evens out the weights, approaching a standard linear regression.



Evaluating Predictive Models

Mean squared error (MSE) or root mean squared error (RMSE) are often used with numeric predictors.

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y - \hat{y})^2$$

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y - \hat{y})^2}$$

RMSE is measured in the same units as the data.

df

	Age	Weight	Pred.Wt	Vax
1	2	25	28.123	0
2	5	37	38.491	1
3	3	33	31.579	0
4	8	44	48.860	1
5	6	50	41.947	0

```
sqrt(mean((df$Weight - df$Pred.Wt)^2))
```

```
[1] 4.5267
```

```
library(caret)
```

```
RMSE(pred=df$Pred.Wt, obs=df$Weight)
```

```
[1] 4.5267
```

Accuracy, Sensitivity and Specificity give information about classifiers.

	Act. Pos	Act. Neg
Pr. Pos	A	B
Pr. Neg	C	D

Accuracy Percentage of correct predictions. $\left(\frac{A+D}{A+B+C+D}\right)$

Sensitivity Percentage of actual positives correctly predicted to be positive. $\left(\frac{A}{A+C}\right)$

Specificity Percentage of actual negatives correctly predicted to be negative. $\left(\frac{D}{B+D}\right)$

```
df$Prob.Vax <-  
  predict(vax.fit, type="response")  
df$Pred.Vax <-  
  as.numeric(df$Prob.Vax > 0.50)  
# Accuracy  
mean(df$Vax == df$Pred.Vax)  
  
[1] 0.6  
# Sensitivity  
with(df[df$Vax==1, ],  
  mean(Vax == Pred.Vax))  
  
[1] 0.5  
# Specificity  
with(df[df$Vax==0, ],  
  mean(Vax == Pred.Vax))  
  
[1] 0.66667
```


For classifiers that use a probability cut-off, the ROC curve describes the trade-off between sensitivity and specificity.

- For a set of thresholds between 0 and 1, predict the binary response.
- Calculate the pair (Specificity, Sensitivity).
- Plot the point on the graph.
- When $AUC \approx 1$, a threshold can be chosen to make both sensitivity and specificity high.

	Vax	Prob.Vax
1	0	0.038597
2	0	0.090001
3	1	0.375094
4	0	0.596565
5	1	0.899743

```
library(pROC)
```

```
vax.roc <- roc(Vax ~ Prob.Vax, data=df,  
plot(vax.roc, type="b"))
```



```
auc(vax.roc)
```

Area under the curve: 0.833

Predictive models should be evaluated on data not used to train the model.

Simplest: Create training and testing sets.

*# Base R: The `sample` command
selects random indices.*

```
ind <- sample(1:nrow(df),  
             round(0.70*nrow(df)))  
df.train <- df[ind, ]  
df.test <- df[-ind, ]  
df.train[, 1:4]
```

	Age	Weight	Pred.Wt	Vax
3	3	33	31.579	0
5	6	50	41.947	0
1	2	25	28.123	0
4	8	44	48.860	1

```
df.test[, 1:4]
```

	Age	Weight	Pred.Wt	Vax
2	5	37	38.491	1

*# Caret's `createDataPartition`
creates indices as well.
Given a factor, it will try to
split them evenly.*

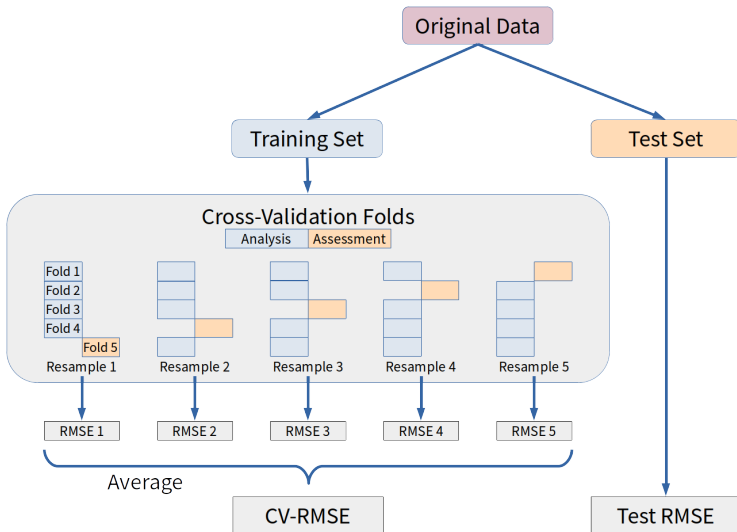
```
ind <- createDataPartition(df$Vax,  
                           p=0.7, list=FALSE)  
df.train <- df[ind, ]  
df.test <- df[-ind, ]  
df.train[, 1:4]
```

	Age	Weight	Pred.Wt	Vax
1	2	25	28.123	0
2	5	37	38.491	1
4	8	44	48.860	1
5	6	50	41.947	0

```
df.test[, 1:4]
```

	Age	Weight	Pred.Wt	Vax
3	3	33	31.579	0

Cross-validation creates many “analysis” and “assessment” sets from the same data set.





Review of R Programming Techniques

Putting commonly-used code in a function can speed coding and prevent copy/paste mistakes.

```
funct.name <- function(<arguments>){...code...}
```

- Multiple commands are placed within braces { }
- Default values can be defined by setting function arguments equal to a value within the definition.
- The return keyword specifies the function return value. Otherwise the last value calculated in the function.
- The warning and stop functions can be used to throw warning and error messages.

```
sumn <- function(x, n=length(x)){sum(x[1:n], na.rm=TRUE)}  
sumn(1:5, n=4)
```

```
[1] 10
```

If dealing with outside data/user input, it's good form to verify inputs. With your own code, it could help debug.

```
sumn <- function(x, n=length(x)){  
  if(!is.numeric(x) || length(dim(x))  
    stop("Requires one-dimensional  
         numeric vector.")  
  if(!is.numeric(n) || length(n) > 1)  
    stop("n must be a single  
         numeric value.")  
  if(round(n) != n){  
    n <- round(n)  
    warning("n must be an integer  
            and will be rounded.")  
  }  
  if(n < 1 || n > length(x))  
    stop("n must be an integer between  
         1 and length of x.")  
  return(sum(x[1:n], na.rm=TRUE))  
}
```

```
a <- 1:5  
b <- c(1, 2, NA, 4, 5)  
sumn(a)
```

```
[1] 15
```

```
sumn(a, n=3)
```

```
[1] 6
```

```
sumn(b, n=3)
```

```
[1] 3
```

```
sumn(a, n=3.2)
```

```
Warning in sumn(a, n = 3.2): n must be  
and will be rounded.
```

```
[1] 6
```

Vectorized functions are usually more efficient than loops.

Many R functions are vectorized.

Square roots of the whole numbers from 1 to 5.

```
a <- 1:5
a
[1] 1 2 3 4 5
sqrt(a)
[1] 1.0000 1.4142 1.7321 2.0000 2.2361
```

Single values are repeated as needed in vectorized arithmetic.

```
# Note: This is a long way to
# calculate variance.
a
[1] 1 2 3 4 5
a.mean <- mean(a)
a.mean
[1] 3
a - a.mean
[1] -2 -1 0 1 2
sum((a - a.mean)^2) / (length(a) - 1)
[1] 2.5
```

The `s/lapply` family of functions iterate functions over vector input. Usually more efficient than a loop.

- `lapply` always returns a list.
- `sapply` tries to return a simplified object.

```
a <- 1:5  
lapply(a, FUN=sqrt)
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] 1.4142
```

```
[[3]]
```

```
[1] 1.7321
```

```
[[4]]
```

```
[1] 2
```

```
sapply(a, FUN=sqrt)
```

```
[1] 1.0000 1.4142 1.7321 2.0000 2.2361
```

Applied to an anonymous function.

```
sapply(a, FUN=function(x){sqrt(x)+2})
```

```
[1] 3.0000 3.4142 3.7321 4.0000 4.2361
```

The `purrr` package (in `tidyverse`) gives the `map` family of commands, which are similar to `s/lapply`.

```
map(1:5, sqrt)
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] 1.4142
```

```
[[3]]
```

```
[1] 1.7321
```

```
[[4]]
```

```
[1] 2
```

```
[[5]]
```

```
[1] 2.2361
```

*# Note that the "suffix" gives the
desired type of output.*

```
map_dbl(1:5, sqrt)
```

```
[1] 1.0000 1.4142 1.7321 2.0000 2.2361
```

*# The "formula" format specifies
an anonymous function compactly.
using `.x` as the input.*

```
map_dbl(1:5, ~ sqrt(.x) + 2)
```

```
[1] 3.0000 3.4142 3.7321 4.0000 4.2361
```

Sometimes a loop is just easier, especially if the function is complicated. So don't feel bad for using one if you have to!

```
for(i in 1:5){  
  # Note that `print` is needed  
  # inside a loop to show results.  
  print(sqrt(i))  
}
```

```
[1] 1  
[1] 1.4142  
[1] 1.7321  
[1] 2  
[1] 2.2361
```

```
# To store loop results, it's good to  
# initialize variables first.  
a.sqrt <- numeric(5)  
# Note that these "type" commands with  
# a number input give empty vectors.
```

```
a.sqrt
```

```
[1] 0 0 0 0 0
```

```
for(i in 1:5){  
  a.sqrt[i] <- sqrt(i)  
}  
a.sqrt
```

```
[1] 1.0000 1.4142 1.7321 2.0000 2.2361
```